

## Mechanistic Evidence

Multiple analyses validate our findings:

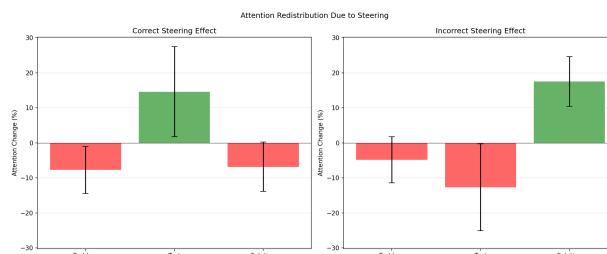
### Prediction Analysis

- Incorrect-predicting: **F1 = 0.821**
- Correct-predicting: F1 = 0.504 (weak)

### Steering Interventions

- Correct-steering fixes 4.04% of errors
- Trade-off: affects 14.66% correct code

### Attention Analysis



Test cases matter MORE than problem descriptions

### Necessity (Orthogonalization)

- Removing correct directions: 83.62% code corruption
- Control removal: only 18.97% corruption

### Persistence Across Fine-tuning

- Base → Instruction-tuned: F1 0.821 → 0.772
- Mechanisms learned in pre-training persist

## Significance

**First application** of Sparse Autoencoders to study code correctness mechanisms in LLMs.

### Practical Applications:

1. **Prompting strategies:** Prioritize test examples over problem descriptions
2. **Error alarms:** Predictor directions flag code for developer review
3. **Selective steering:** Intervene only when errors anticipated

**Safety Implications:** Contributes to safer AI deployment in healthcare, finance, and critical infrastructure.

## Key References

- Bricken et al. (2023) — Towards Monosemanticity
- Templeton et al. (2024) — Scaling Monosemanticity
- Lieberum et al. (2024) — GemmaScope
- Ferrando et al. (2024) — Entity Recognition via SAEs

## Contact

### Proponent:

Kriz Royce Tahimic  
[kriz\\_tahimic@dlsu.edu.ph](mailto:kriz_tahimic@dlsu.edu.ph)

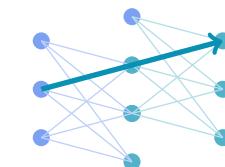
### Adviser:

Dr. Charibeth K. Cheng  
[charibeth.cheng@dlsu.edu.ph](mailto:charibeth.cheng@dlsu.edu.ph)

College of Computer Studies  
De La Salle University  
Manila, Philippines

# Mechanistic Interpretability of Code Correctness in LLMs

via Sparse Autoencoders



**Kriz Royce Tahimic**

Adviser: Dr. Charibeth K. Cheng

College of Computer Studies  
De La Salle University

Academic Year 2024–2025

## Context

AI-assisted coding has reached critical mass:

- 30% of GitHub Copilot suggestions enter production code
- Projected **\$1.5 trillion** GDP impact by 2030
- LLMs increasingly deployed in critical systems

## The Problem

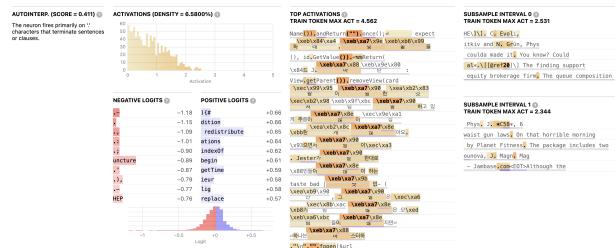
We lack **mechanistic understanding** of WHEN and WHY models produce correct code.

- Models fail in bug-prone contexts (12.27% accuracy)
- 44% of LLM bugs mirror historical training bugs
- Critical for: healthcare, banking, military

## The Challenge

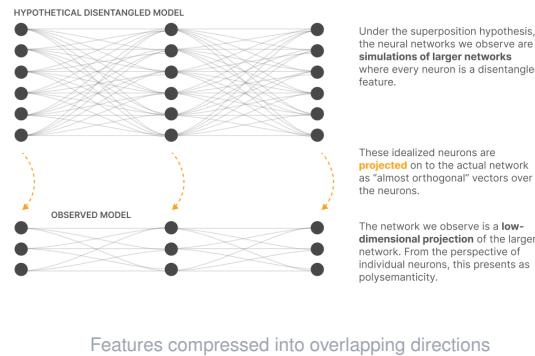
Neural networks are hard to interpret:

- **Polysemantic neurons:** One neuron responds to multiple unrelated concepts



## Why Interpretation is Hard

**Superposition:** Networks compress more features than dimensions, entangling representations.



## Our Approach

**Sparse Autoencoders (SAEs)** decompose entangled representations into interpretable directions.

**Key Idea:** If code correctness is represented as a *direction* in the model's latent space, we can:

1. **Detect** it (prediction)
2. **Manipulate** it (steering)
3. **Validate** it (causal analysis)

### Technical Setup:

- **Model:** Gemma-2 2B (base & instruction-tuned)
- **SAE:** GemmaScope pre-trained (16K features)
- **Dataset:** MBPP (1,000 Python problems)
- **Analysis:** Residual stream at layer 20

## Key Discovery

**Code correctness directions EXIST** in LLM representations—and they are **actionable**.

### 1. Predict Errors Before Generation

The **incorrect-predicting direction** detects errors with high accuracy:

#### POSITIVE LOGITS

none	1.35
None	1.28
none	1.24
None	1.21
SourceChecksum	1.01
NONE	0.96
NONE	0.94
autorytatywna	0.89
なし	0.87
لادسما	0.83

F1: 0.821 for error detection—can serve as “error alarm”

### 2. Steer Toward Correctness

The **correct-steering direction** can fix errors (4.04% of incorrect code fixed).

### 3. Asymmetric Finding

- Found: **incorrect-predicting + correct-steering**
- Did NOT find: correct-predicting or incorrect-steering
- *Models detect “wrongness” differently than they encode “correctness”*