

## ◆ Context

AI-assisted coding has reached **critical mass**:

- 30% of GitHub Copilot suggestions enter production
  - Projected **\$1.5T** GDP impact by 2030
  - LLMs deployed in critical systems worldwide

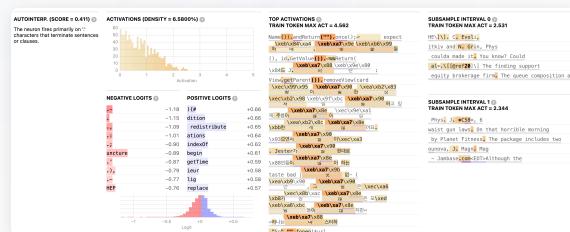
## ◆ The Problem

We lack **mechanistic understanding** of WHEN and WHY models produce correct code.

- Models fail in bug-prone contexts (12.27% accuracy)
  - 44% of LLM bugs mirror historical training bugs
  - **Stakes:** Healthcare, banking, military

## ◆ The Challenge

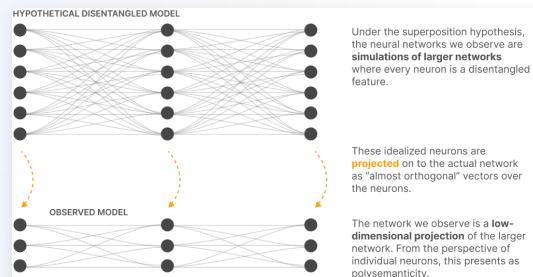
**Polysemantic neurons:** One neuron responds to multiple unrelated concepts



A single neuron activates for unrelated concepts

## ◆ Why Interpretation is Hard

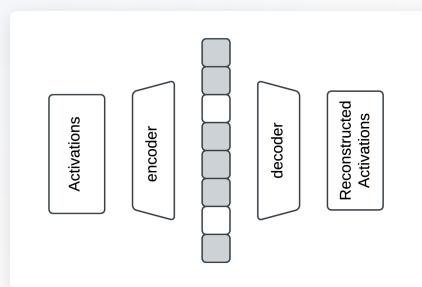
**Superposition:** Networks compress more features than dimensions, entangling representations.



*Observed models are low-dimensional projections of larger networks*

## ◆ Our Approach

**Sparse Autoencoders (SAEs)** decompose entangled representations into interpretable directions.



SAE: Activations  $\rightarrow$  sparse latent space  $\rightarrow$  reconstructed activations

## Technical Setup

<b>Model</b>	Gemma-2 2B	<b>SAE</b>	GemmaScope (16K latents)
<b>Dataset</b>	MBPP (1K problems)	<b>Analysis</b>	All layers, residual stream, final prompt token

## ◆ Key Discovery

**Code correctness directions EXIST in LLM representations—and they are actionable.**

## 1. Predict Errors Before Generation

The **incorrect-predicting direction** detects errors with high accuracy:

POSITIVE LOGITS	①
none	1.35
None	1.28
none	1.24
None	1.21
SourceChecksum	1.01
NONE	0.96
NONE	0.94
autorytatywna	0.89
なし	0.87
لاصقا	0.83

E1: 0.821 for error detection—can serve as "error alarm"

## 2. Steer Toward Correctness

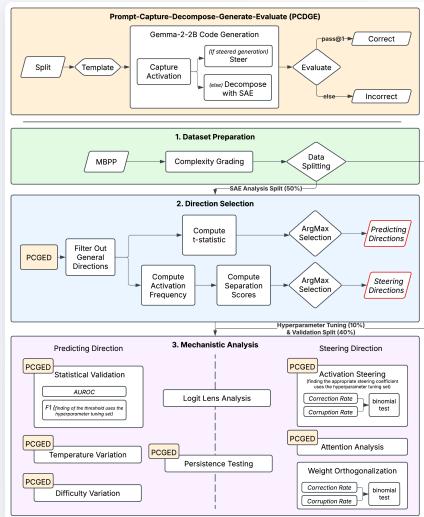
The **correct-steering direction** can fix errors (4.04% of incorrect code fixed).

### 3. Asymmetric Finding

**Found:** incorrect-predicting + correct-steering  
**Not Found:** correct-predicting or incorrect-steering

*Models detect "wrongness" differently than they encode "correctness"*

## ◆ Mechanistic Evidence



PCDGE pipeline: Dataset preparation → Direction selection → Mechanistic analysis

### Prediction Analysis

- Incorrect-predicting:  $F1 = 0.821$
- Correct-predicting:  $F1 = 0.504$  (weak)
- Correct-steering fixes **4.04%** of errors
- Trade-off: affects 14.66% correct code

**Causal validation:** Removing directions causes 83.62% corruption (vs. 18.97% control). Directions persist across instruction-tuning ( $F1: 0.821 \rightarrow 0.772$ ).

## ◆ Significance

First application of Sparse Autoencoders to study code correctness mechanisms in LLMs.

### Practical Applications

- Prompting strategies:** Prioritize test examples over problem descriptions
- Error alarms:** Predictor directions flag code for review
- Selective steering:** Intervene only when errors anticipated

**Safety Implications:** Contributes to safer AI deployment in healthcare, finance, and critical infrastructure.

### Key References

- Bricken et al. (2023) — Towards Monosemantics  
Templeton et al. (2024) — Scaling Monosemantics  
Lieberum et al. (2024) — GemmaScope  
Ferrando et al. (2024) — Entity Recognition via SAEs

### Contact

**PROONENT**  
Kriz Tahimic  
kriz\_tahimic@dlsu.edu.ph

**ADVISER**  
Dr. Charibeth K. Cheng  
charibeth.cheng@dlsu.edu.ph

College of Computer Studies  
De La Salle University, Manila  
BS Computer Science — Software Technology

# Mechanistic Interpretability

## of Code Correctness in LLMs

via Sparse Autoencoders



Kriz Tahimic

Adviser: Dr. Charibeth K. Cheng

College of Computer Studies  
De La Salle University

Academic Year 2024–2025